

S&DS 365 / 565
Data Mining and Machine Learning

Yale

For today

- Finding solutions with generative models
- Recall linear regression solution
- iterative methods
- optimization error versus statistical error

Recap

Exponential families: power probabilistic model

Can use to **roughly** approximate $\mathbb{P}_\beta(y|x)$

Classification: $\hat{f}(x) = \mathbb{1}\left[\mathbb{P}_\beta(y = 1|x) > \mathbb{P}_\beta(y = 0|x)\right]$

How do we find parameters?

Start with different probabilistic model

Generative modeling

Exponential families focus on $\mathbb{P}(y|x)$

Generative modeling: $\mathbb{P}(x|y)$

Want $\mathbb{P}(y|x)$

Generative modeling

Bayes rule

$$\mathbb{P}(y|x) = \frac{\overset{\mathbb{P}}{\cancel{\text{prob}}}(x|y)\mathbb{P}(y)}{\mathbb{P}(x)}$$

Classification:

$$\mathbb{P}(y) = \begin{cases} \pi_0 & y = 0 \\ \pi_1 & y = 1 \end{cases}$$

$$\mathbb{P}(x) = \sum_j \mathbb{P}(x|y = j)\mathbb{P}(y = j)$$

Generative modeling

Modeling choice: $\mathbb{P}(x|y)$

Gaussian: $x|y \sim N(\mu_y, \Sigma_y)$ (exploring in homework)

Naive Bayes: $x \in \mathbb{R}^d$ with

$$\mathbb{P}(x|y) = \prod_{j=1}^d \mathbb{P}(x_{(j)}|y)$$

Each coordinate is **independent** conditioned on y

Naive Bayes

Can pick any univariate distribution for $x_i|y$

Binary classification:

$$\mathbb{P}(y = 1|x) > \mathbb{P}(y = 0|x)$$

Naive Bayes

Can pick any univariate distribution for $x_i|y$

Binary classification:

$$\mathbb{P}(y = 1|x) > \mathbb{P}(y = 0|x)$$

$$\frac{\mathbb{P}(x|y = 1)\mathbb{P}(y = 1)}{\mathbb{P}(x)} > \frac{\mathbb{P}(x|y = 0)\mathbb{P}(y = 0)}{\mathbb{P}(x)}$$

Naive Bayes

Can pick any univariate distribution for $x_i|y$

Binary classification:

$$\mathbb{P}(y = 1|x) > \mathbb{P}(y = 0|x)$$

$$\frac{\mathbb{P}(x|y = 1)\mathbb{P}(y = 1)}{\mathbb{P}(x)} > \frac{\mathbb{P}(x|y = 0)\mathbb{P}(y = 0)}{\mathbb{P}(x)}$$

$$\mathbb{P}(x|y = 1)\mathbb{P}(y = 1) > \mathbb{P}(x|y = 0)\mathbb{P}(y = 0)$$

Naive Bayes

Can pick any univariate distribution for $x_i|y$

Binary classification:

$$\mathbb{P}(y = 1|x) > \mathbb{P}(y = 0|x)$$

$$\frac{\mathbb{P}(x|y = 1)\mathbb{P}(y = 1)}{\mathbb{P}(x)} > \frac{\mathbb{P}(x|y = 0)\mathbb{P}(y = 0)}{\mathbb{P}(x)}$$

$$\mathbb{P}(x|y = 1)\mathbb{P}(y = 1) > \mathbb{P}(x|y = 0)\mathbb{P}(y = 0)$$

$$\log(\mathbb{P}(x|y = 1)) + \log(\mathbb{P}(y = 1)) > \log(\mathbb{P}(x|y = 0)) + \log(\mathbb{P}(y = 0))$$

Naive Bayes

Can pick any univariate distribution for $x_i|y$

Binary classification:

$$\mathbb{P}(y = 1|x) > \mathbb{P}(y = 0|x)$$

$$\frac{\mathbb{P}(x|y = 1)\mathbb{P}(y = 1)}{\mathbb{P}(x)} > \frac{\mathbb{P}(x|y = 0)\mathbb{P}(y = 0)}{\mathbb{P}(x)}$$

$$\mathbb{P}(x|y = 1)\mathbb{P}(y = 1) > \mathbb{P}(x|y = 0)\mathbb{P}(y = 0)$$

$$\log(\mathbb{P}(x|y = 1)) + \log(\mathbb{P}(y = 1)) > \log(\mathbb{P}(x|y = 0)) + \log(\mathbb{P}(y = 0))$$

$$\sum_{j=1}^d \log \left(\frac{\mathbb{P}(x_{(j)}|y = 1)}{\mathbb{P}(x_{(j)}|y = 0)} \right) > \log \left(\frac{\pi_0}{\pi_1} \right)$$

Example: spam

Each feature is count of a word

More easily each feature $x_{(j)} \in \{0, 1\}$ is if word j appears or not

$$\mathbb{P}(x_{(j)}|y) = p_{j,y}^{x_{(j)}}(1 - p_{j,y})^{1-x_{(j)}}$$

Example: spam

Each feature is count of a word

More easily each feature $x_{(j)} \in \{0, 1\}$ is if word j appears or not

$$\mathbb{P}(x_{(j)}|y) = p_{j,y}^{x_{(j)}}(1 - p_{j,y})^{1-x_{(j)}}$$

$$\sum_{j=1}^d \log \left(\frac{p_{j,1}^{x_{(j)}}(1 - p_{j,1})^{1-x_{(j)}}}{p_{j,0}^{x_{(j)}}(1 - p_{j,0})^{1-x_{(j)}}} \right) > \log \left(\frac{\pi_0}{\pi_1} \right)$$

Example: spam

Each feature is count of a word

More easily each feature $x_{(j)} \in \{0, 1\}$ is if word j appears or not

$$\mathbb{P}(x_{(j)}|y) = p_{j,y}^{x_{(j)}}(1 - p_{j,y})^{1-x_{(j)}}$$

$$\sum_{j=1}^d \log \left(\frac{p_{j,1}^{x_{(j)}}(1 - p_{j,1})^{1-x_{(j)}}}{p_{j,0}^{x_{(j)}}(1 - p_{j,0})^{1-x_{(j)}}} \right) > \log \left(\frac{\pi_0}{\pi_1} \right)$$

$$\sum_{j=1}^d x_{(j)} \log \left(\frac{p_{j,1}}{1 - p_{j,1}} \frac{1 - p_{j,0}}{p_{j,0}} \right) + \log \left(\frac{1 - p_{j,1}}{1 - p_{j,0}} \right) > \log \left(\frac{\pi_0}{\pi_1} \right)$$

Linear classifier!

Naive Bayes

Can't actually use in practice. Don't have $\mathbb{P}(x_{(j)}|y)$

Spam example: Estimate $p_{j,y}$ from training data!

Very simple procedure: called plug-in methods

Move back to general optimization

Finding solutions

- For linear regression we have normal equations

$$\min_{\beta} \frac{1}{2} \|X\beta - y\|_2^2 \quad \min_{\beta} \frac{1}{2} \sum_{i=1}^n (x_i^T \beta - y_i)^2$$

$$\nabla \frac{1}{2} \|X\beta - y\|_2^2 = X^T(X\beta - y)$$

set to zero

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

(can use psuedo-inverse if not invertible)

Finding solutions with optimization


- Let $f(\beta)$ be some convex function
- Goal: solve

$$\arg \min_{\beta} f(\beta)$$

- Eg $f(\beta) = \frac{1}{2n} \|X\beta - y\|_2^2$
- Logistic: $f(\beta) = \frac{1}{n} \sum_i -y_i x_i^T \beta + \log(1 + \exp(x_i^T \beta))$

Gradients, descents, and recall Taylor

Taylor: $f(\beta) \approx f(\beta_0) + (\beta - \beta_0)^T \nabla f(\beta_0)$ for β close to β_0

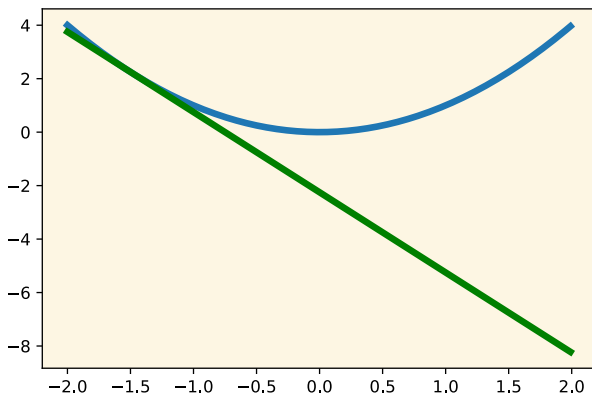


Gradients, descents, and recall Taylor

Taylor: $f(\beta) \approx f(\beta_0) + (\beta - \beta_0)^T \nabla f(\beta_0)$ for β close to β_0

$$\beta = \beta_0 - \eta \nabla f(\beta_0)$$

$$f(\beta) \approx f(\beta_0) - \eta \|\nabla f(\beta_0)\|^2$$



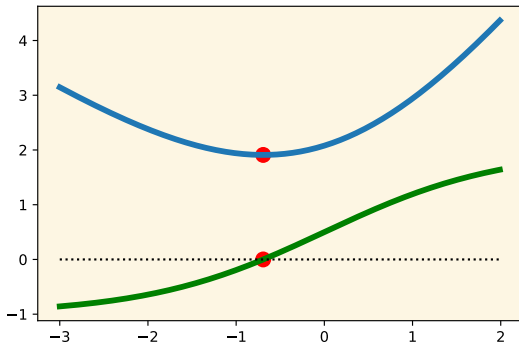
Gradient Descent

$$\beta_k = \beta_{k-1} - \eta_{k-1} \nabla f(\beta_{k-1})$$

Newton and Taylor

Remember trying to set $\nabla f(\beta) = 0$

Consider: $f(\beta) = 3 \log(1 + \exp(\beta)) - \beta$ $= \sum_{i=1}^3 \log(1 + \exp(\beta)) - y_i \beta$
 $y_1 = 1 \quad y_2 = 0 \quad y_3 = 0$




$$\nabla f(\beta) \approx \nabla f(\beta_0) + H f(\beta_0)(\beta - \beta_0)$$

$H f(\beta)$ is the Hessian

Consider: $f(\beta) = 3 \log(1 + \exp(\beta)) - \beta$

$$f'(\beta) = 3 \frac{\exp(\beta)}{1 + \exp(\beta)} - 1$$




$$\nabla f(\beta) \approx \nabla f(\beta_0) + H f(\beta_0)(\beta - \beta_0)$$

$H f(\beta)$ is the Hessian

Consider: $f(\beta) = 3 \log(1 + \exp(\beta)) - \beta$

$$f'(\beta) = 3 \frac{\exp(\beta)}{1 + \exp(\beta)} - 1$$




$$\nabla f(\beta) \approx \nabla f(\beta_0) + H f(\beta_0)(\beta - \beta_0)$$

$H f(\beta)$ is the Hessian

Consider: $f(\beta) = 3 \log(1 + \exp(\beta)) - \beta$

$$f'(\beta) = 3 \frac{\exp(\beta)}{1 + \exp(\beta)} - 1$$




$$\nabla f(\beta) \approx \nabla f(\beta_0) + H f(\beta_0)(\beta - \beta_0)$$

$H f(\beta)$ is the Hessian

Consider: $f(\beta) = 3 \log(1 + \exp(\beta)) - \beta$

$$f'(\beta) = 3 \frac{\exp(\beta)}{1 + \exp(\beta)} - 1$$



Newton and Taylor

- $\nabla f(\beta) \approx \nabla f(\beta_0) + Hf(\beta_0)(\beta - \beta_0) = 0 \Rightarrow \beta = \beta_0 - [Hf(\beta_0)]^{-1} \nabla f(\beta_0)$
- $Hf(\beta)$ is the Hessian

$$[Hf(\beta_k)]_{ij} = \frac{\partial^2 f(\beta_k)}{\partial \beta_i \partial \beta_j}$$

- Newton's method $\beta_k = \beta_{k-1} - (Hf(\beta_{k-1}))^{-1} \nabla f(\beta_{k-1})$
- Computational complexity is $O(p^3)$ in general. For large problems impractical.

Newton's method and Linear Regression

What is it?

$$f(\beta) = \frac{1}{2} \|X\beta - y\|_2^2$$

$$\nabla f(\beta) = X^T (X\beta - y)$$

$$H(f(\beta)) = X^T X$$

$$\beta_k = \beta_{k-1} - (X^T X)^{-1} (X^T (X\beta_{k-1} - y))$$

$$= \beta_{k-1} - (\beta_{k-1} - (X^T X)^{-1} X^T y)$$

$$= (X^T X)^{-1} X^T y$$

$$f(\beta) = \frac{1}{2} \sum_{i=1}^n (x_i^T \beta - y_i)^2 \xrightarrow{\nabla} \frac{1}{2} \sum_{i=1}^n \nabla (x_i^T \beta - y_i)^2 \rightarrow \sum_{i=1}^n (x_i^T \beta - y_i) x_i$$

$$\nabla f(\beta) = \sum_{i=1}^n x_i (x_i^T \beta - y_i) \xrightarrow{\text{Hessian}} \sum_{i=1}^n x_i (x_i^T) = X^T X$$

$$[H f(\beta)]_{(ab)} = \frac{\partial (\nabla f(\beta))_a}{\partial \beta_{(b)}} = \frac{\partial (\sum_{i=1}^n x_{i(a)} (x_i^T \beta - y_i))}{\partial \beta_{(b)}}$$

$$= \sum_{i=1}^n x_{i(a)} \frac{\partial (x_i^T \beta - y_i)}{\partial \beta_{(b)}} = \sum_{i=1}^n x_{i(a)} x_{i(b)}$$

Gradient descent

- Used extensively for large scale problems
- $\beta_k = \beta_{k-1} - \eta_{k-1} \nabla f(\beta_{k-1})$
- but “slower” than Newton
- Before looking at figures, what is the update for linear regression?

Gradient descent and Linear Regression

Recall that $f(\beta) = \frac{1}{2n} \|X\beta - y\|_2^2$

$$\nabla f(\beta) = \frac{1}{n} X^T (X\beta - y)$$

$$\beta_k = \beta_{k-1} - \eta_{k-1} \cdot \frac{1}{n} (X^T (X\beta_{k-1} - y))$$

Gradient descent and Linear Regression (2)

Recall that $f(\beta) = \frac{1}{2n} \sum_{i=1}^n (x_i^T \beta - y_i)^2$

$$\beta_k = \beta_{k-1} - \frac{n_{k-1}}{n} \sum_{i=1}^n x_i \underbrace{(x_i^T \beta_{k-1} - y_i)}_{\text{error}_i}$$

Exponential families:

$$f(B) = \frac{1}{n} \sum_{i=1}^n A(x_i^T B) - y_i x_i^T B$$

recall

$$A: \mathbb{R} \rightarrow \mathbb{R}$$

$$\nabla f(B) = \frac{1}{n} \sum_{i=1}^n A'(x_i^T B) x_i - y_i x_i$$

$$= \frac{1}{n} \sum_{i=1}^n \underbrace{(A'(x_i^T B) - y_i)}_{\text{error}_i} x_i$$

$$Hf(B) = \frac{1}{n} \sum_{i=1}^n A''(x_i^T B) x_i x_i^T$$

$$\text{error} \in \mathbb{R}^n$$

$$\text{error}_i = A'(x_i^T B) - y_i$$

$$\text{Linear: } A(\theta) = \frac{1}{2} \theta^2$$

$$\text{Logistic: } A(\theta) = \log(1 + \exp(\theta))$$

$$A'(\theta) = \frac{\exp(\theta)}{1 + \exp(\theta)}$$

Stochastic Gradient Descent

- gradient descent cycles through all of the data
- large n , problematic
- stochastic gradient descent (sgd)

Stochastic Gradient Descent

- Recall that $f(\beta) = \frac{1}{2n} \sum_{i=1}^n (x_i^T \beta - y_i)^2$
- gradient descent is

$$\beta_k = \beta_{k-1} - \eta_{k-1} \nabla f(\beta_{k-1})$$

- For linear regression that becomes

$$\beta_k = \beta_{k-1} - \eta_{k-1} \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i (\mathbf{x}_i^T \beta - y_i)$$

- Let's be lazy and not go through all of the data

Stochastic Gradient Descent

- Recall that $f(\beta) = \frac{1}{2n} \sum_{i=1}^n (x_i^T \beta - y_i)^2$
- gradient descent is

$$\beta_k = \beta_{k-1} - \eta_{k-1} \nabla f(\beta_{k-1})$$

- For linear regression that becomes

$$\beta_k = \beta_{k-1} - \eta_{k-1} \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i (\mathbf{x}_i^T \beta - y_i)$$

- Let's be lazy and not go through all of the data

$$\beta_k = \beta_{k-1} - \eta_{k-1} \mathbf{x}_J (\mathbf{x}_J^T \beta_{k-1} - y_J)$$

- $J \in [n]$ is uniform random variable
- iterative learning algo applying stochastic gradient descent (sgd)

Stochastic Gradient Descent

- Recall that $f(\beta) = \frac{1}{2n} \sum_{i=1}^n (x_i^T \beta - y)^2$
- We can write $f(\beta) = \frac{1}{n} \sum_{i=1}^n f_i(\beta)$

$$f_i(\beta) = \frac{1}{2} (x_i^T \beta - y)^2$$

- So batch gradient descent

$$\beta_k = \beta_{k-1} - \eta_{k-1} \frac{1}{n} \sum_{i=1}^n \nabla f_i(\beta_{k-1})$$

- stochastic gradient descent (sgd) is

$$\beta_k = \beta_{k-1} - \eta_{k-1} \nabla f_j(\beta_{k-1})$$

Stochastic Gradient Descent

- Most general version of sgd
- Goal: $\arg \min_{\beta} f(\beta)$
- Have access to random vector $\mathbb{E} g_{\beta} = \nabla f(\beta)$


$$\beta_k = \beta_{k-1} - \eta_{k-1} g_{\beta_{k-1}}$$

random!


- In the previous slides randomness over J the random example

$$\begin{aligned}\mathbb{E} \nabla f_J(\beta) &= \frac{1}{n} \sum_{j=1}^n \nabla (x_j^T \beta - y_j)^2 = \frac{1}{n} \sum_{j=1}^n \nabla f_j(\beta) \\ &= \nabla f(\beta)\end{aligned}$$

Convergence grad v newton




Convergence grad v newton against time



Convergence grad v newton error to β^*


Notion of statistical error versus optimization error

$$y = X\beta^* + w$$

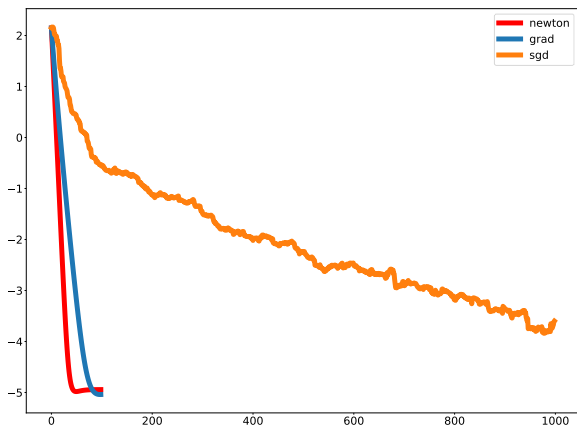


Convergence grad v newton error to β^*

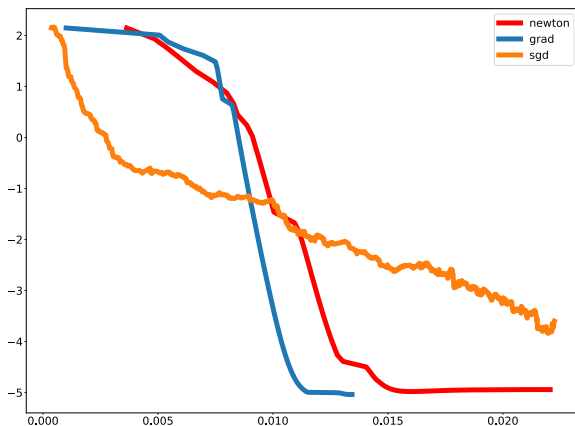
Notion of statistical error versus optimization error vs time



Convergence grad v newton v sgd error to β^*



Convergence grad v newton v sgd error to β^*



Move to notebook